

```
!git clone https://bitbucket.org/jadslim/german-traffic-signs
```

```
Cloning into 'german-traffic-signs'...  
Unpacking objects: 100% (6/6), done.
```

```
import keras  
import pickle  
import random  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from keras.models import Sequential  
from keras.layers import Dense  
from tensorflow.keras.optimizers import Adam  
from keras.utils.np_utils import to_categorical  
from keras.layers import Dropout, Flatten  
from keras.layers.convolutional import Conv2D, MaxPooling2D
```

```
np.random.seed(0)
```

```
with open('german-traffic-signs/train.p','rb') as f:  
    train_data = pickle.load(f)  
with open('german-traffic-signs/valid.p','rb') as f:  
    valid_data = pickle.load(f)  
with open('german-traffic-signs/test.p','rb') as f:  
    test_data = pickle.load(f)
```

```
X_train,y_train = train_data['features'],train_data['labels']  
X_valid,y_valid = valid_data['features'],valid_data['labels']  
X_test,y_test = test_data['features'],test_data['labels']
```

```
print(X_train.shape)  
print(X_valid.shape)  
print(X_test.shape)
```

```
(34799, 32, 32, 3)  
(4410, 32, 32, 3)  
(12630, 32, 32, 3)
```

```
assert(X_train.shape[0]==y_train.shape[0]),"The number of images is not equal to the number of labels"
assert(X_valid.shape[0]==y_valid.shape[0]),"The number of images is not equal to the number of labels"
assert(X_test.shape[0]==y_test.shape[0]),"The number of images is not equal to the number of labels"
assert(X_train.shape[1:]==(32,32,3)), "The dimensions of the images are not 32 x 32 x 3"
assert(X_valid.shape[1:]==(32,32,3)), "The dimensions of the images are not 32 x 32 x 3"
assert(X_test.shape[1:]==(32,32,3)), "The dimensions of the images are not 32 x 32 x 3"
```

```
data = pd.read_csv('german-traffic-signs/signnames.csv')
```

```
num_of_samples=[]
```

```
cols = 5
```


```
num_classes = 43
```

```
fig, axs = plt.subplots(nrows=num_classes, ncols=cols, figsize=(5,50))
fig.tight_layout()
```


```
for i in range(cols):
    for j,row in data.iterrows():
        x_selected = X_train[y_train == j]
        axs[j][i].imshow(x_selected[random.randint(0,(len(x_selected) - 1)), :, :], cmap=plt.get_cmap('gray'))
        axs[j][i].axis("off")
        if i == 2:
            axs[j][i].set_title(str(j) + "-" + row["SignName"])
            num_of_samples.append(len(x_selected))
```



0-Speed limit (20km/h)




1-Speed limit (30km/h)



2-Speed limit (50km/h)




3-Speed limit (60km/h)




4-Speed limit (70km/h)




5-Speed limit (80km/h)




6-End of speed limit (80km/h)



7-Speed limit (100km/h)



8-Speed limit (120km/h)




9-No passing






10-No passing for vehicles over 3.5 metric tons



11-Right-of-way at the next intersection



12-Priority road




13-Yield



14-Stop




15-No vehicles




16-Vehicles over 3.5 metric tons prohibited



17-No entry



18-General caution




19-Dangerous curve to the left



20-Dangerous curve to the right



21-Double curve



22-Bumpy road



23-Slippery road



24-Road narrows on the right




25-Road work



26-Traffic signals



27-Pedestrians



28-Children crossing





29-Bicycles crossing




30-Beware of ice/snow




31-Wild animals crossing



32-End of all speed and passing limits



33-Turn right ahead




34-Turn left ahead




35-Ahead only



36-Go straight or right




37-Go straight or left





```
print(num_of_samples)
plt.figure(figsize=(12, 4))
plt.bar(range(0, num_classes), num_of_samples)
plt.title("Distribution of the train dataset")
plt.xlabel("Class number")
plt.ylabel("Number of images")
plt.show()
```

[180, 1980, 2010, 1260, 1770, 1650, 360, 1290, 1260, 1320, 1800, 1170, 1890, 1920, 690, 540, 360, 990, 1080, 180, 300, 270, 330, 450, 240]




```
import cv2

plt.imshow(X_train[1000])
plt.axis("off")
print(X_train[1000].shape)
print(y_train[1000])
```

(32, 32, 3)


36



```
def grayscale(img):  
    img = cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)  
    return img
```

```
img = grayscale(X_train[1000])  
plt.imshow(img)  
plt.axis('off')  
print(img.shape)
```


(32, 32)




```
def equalize(img):  
    img = cv2.equalizeHist(img)  
    return img
```

```
img = equalize(img)  
plt.imshow(img)  
plt.axis('off')  
print(img.shape)
```

(32, 32)




```
def preprocessing(img):  
    img = grayscale(img)  
    img = equalize(img)  
    img = img/255  
    return img
```

```
X_train = np.array(list(map(preprocessing,X_train)))  
X_valid = np.array(list(map(preprocessing,X_valid)))  
X_test = np.array(list(map(preprocessing,X_test)))
```

```
plt.imshow(X_train[random.randint(0,len(X_train)-1)])  
plt.axis('off')  
print(X_train.shape)
```

(34799, 32, 32)



```
X_train = X_train.reshape(34799,32,32,1)
X_test = X_test.reshape(12630,32,32,1)
X_valid = X_valid.reshape(4410,32,32,1)
```

```
from keras.preprocessing.image import ImageDataGenerator
```


```
datagen = ImageDataGenerator(width_shift_range=0.1,
                              height_shift_range=0.1,
                              zoom_range=0.2,
                              shear_range=0.1,
                              rotation_range=10.)
```

```
datagen.fit(X_train)
```

```
batches = datagen.flow(X_train,y_train,batch_size=15)
X_batch,y_batch = next(batches)
```

```
fig,axs = plt.subplots(1,15,figsize=(20,5))
fig.tight_layout()
```

```
for i in range(15):
    axs[i].imshow(X_batch[i].reshape(32,32))
    axs[i].axis('off')
```



```
print(X_batch.shape)
```

```
(15, 32, 32, 1)
```

```
y_train = to_categorical(y_train,43)
y_test = to_categorical(y_test,43)
y_valid = to_categorical(y_valid,43)
```

```
def modified_model():
    model = Sequential()

    model.add(Conv2D(60,(5,5),input_shape=(32,32,1),activation='relu'))
    model.add(Conv2D(60,(5,5),activation='relu'))
    model.add(MaxPooling2D(pool_size=(2,2)))

    model.add(Conv2D(30,(3,3),activation='relu'))
    model.add(Conv2D(30,(3,3),activation='relu'))
    model.add(MaxPooling2D(pool_size=(2,2)))

    model.add(Flatten())
    model.add(Dense(500,activation='relu'))
    # model.add(Dropout(0.5))
    model.add(Dense(num_classes,activation='softmax'))

    model.compile(Adam(learning_rate=0.001),loss='categorical_crossentropy',metrics=['accuracy'])
    return model
```

```
model = modified_model()
print(model.summary())
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 28, 28, 60)	1560
conv2d_9 (Conv2D)	(None, 24, 24, 60)	90060
max_pooling2d_4 (MaxPooling	(None, 12, 12, 60)	0

```

2D)

conv2d_10 (Conv2D)          (None, 10, 10, 30)      16230

conv2d_11 (Conv2D)          (None, 8, 8, 30)        8130

max_pooling2d_5 (MaxPooling (None, 4, 4, 30)      0
2D)

flatten_2 (Flatten)         (None, 480)             0

dense_4 (Dense)             (None, 500)             240500

dense_5 (Dense)             (None, 43)              21543

```

```

=====
Total params: 378,023
Trainable params: 378,023
Non-trainable params: 0

```

None

```

history = model.fit(datagen.flow(X_train,y_train,batch_size=50),epochs=10,validation_data=(X_valid, y_valid),verbose=1,shuffle=1)

```


```

Epoch 1/10
696/696 [=====] - 14s 19ms/step - loss: 1.3104 - accuracy: 0.6236 - val_loss: 0.1919 - val_accuracy: 0.9429
Epoch 2/10
696/696 [=====] - 13s 19ms/step - loss: 0.3413 - accuracy: 0.8949 - val_loss: 0.0985 - val_accuracy: 0.9685
Epoch 3/10
696/696 [=====] - 13s 19ms/step - loss: 0.2150 - accuracy: 0.9338 - val_loss: 0.0922 - val_accuracy: 0.9732
Epoch 4/10
696/696 [=====] - 14s 20ms/step - loss: 0.1582 - accuracy: 0.9511 - val_loss: 0.0734 - val_accuracy: 0.9780
Epoch 5/10
696/696 [=====] - 14s 20ms/step - loss: 0.1273 - accuracy: 0.9599 - val_loss: 0.0849 - val_accuracy: 0.9748
Epoch 6/10
696/696 [=====] - 13s 19ms/step - loss: 0.1129 - accuracy: 0.9646 - val_loss: 0.0755 - val_accuracy: 0.9771
Epoch 7/10
696/696 [=====] - 14s 20ms/step - loss: 0.0980 - accuracy: 0.9694 - val_loss: 0.0576 - val_accuracy: 0.9832
Epoch 8/10
696/696 [=====] - 13s 19ms/step - loss: 0.0877 - accuracy: 0.9725 - val_loss: 0.0511 - val_accuracy: 0.9844
Epoch 9/10
696/696 [=====] - 13s 19ms/step - loss: 0.0794 - accuracy: 0.9751 - val_loss: 0.0451 - val_accuracy: 0.9880
Epoch 10/10
696/696 [=====] - 14s 20ms/step - loss: 0.0778 - accuracy: 0.9757 - val_loss: 0.0469 - val_accuracy: 0.9889

```


```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.legend(['training', 'validation'])
plt.title('Loss')
plt.xlabel('epoch')
```

Text(0.5, 0, 'epoch')



```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'validation'])
plt.title('Accuracy')
plt.xlabel('epoch')
```

```
Text(0.5, 0, 'epoch')
```



```
score = model.evaluate(X_test,y_test,verbose=0)
print('Test Score',score[0])
print('Test Accuracy',score[1])
```

```
Test Score 0.15319083631038666
Test Accuracy 0.9634996056556702
```

acc | / | training ||

```
import requests
from PIL import Image
url1 = 'https://c8.alamy.com/comp/G667W0/road-sign-speed-limit-30-kmh-zone-passau-bavaria-germany-G667W0.jpg' #1
url2 = 'https://c8.alamy.com/comp/A0RX23/cars-and-automobiles-must-turn-left-ahead-sign-A0RX23.jpg' #34
url3 = 'https://previews.123rf.com/images/bwylezich/bwylezich1608/bwylezich160800375/64914157-german-road-sign-slippery-road.jpg' #23
url4 = 'https://previews.123rf.com/images/pejo/pejo0907/pejo090700003/5155701-german-traffic-sign-no-205-give-way.jpg' #13
url5 = 'https://c8.alamy.com/comp/J2MRAJ/german-road-sign-bicycles-crossing-J2MRAJ.jpg' #29
```

```
r = requests.get(url1, stream=True)
img = Image.open(r.raw)
plt.imshow(img, cmap=plt.get_cmap('gray'))
img = np.asarray(img)
img = cv2.resize(img, (32, 32))
img = preprocessing(img)
plt.imshow(img, cmap = plt.get_cmap('gray'))
print(img.shape)
```

(32, 32)




```
img = img.reshape(1, 32, 32, 1)
print("Predicted sign: "+ str(np.argmax(model.predict(img), axis=-1)))
```

Predicted sign: [1]



```
r = requests.get(url2, stream=True)
img = Image.open(r.raw)
plt.imshow(img, cmap=plt.get_cmap('gray'))
img = np.asarray(img)
img = cv2.resize(img, (32, 32))
img = preprocessing(img)
plt.imshow(img, cmap = plt.get_cmap('gray'))
print(img.shape)
```

(32, 32)




```
img = img.reshape(1, 32, 32, 1)
print("Predicted sign: "+ str(np.argmax(model.predict(img), axis=-1)))
```

Predicted sign: [34]

```
r = requests.get(url3, stream=True)
```

```
img = Image.open(r.raw)
plt.imshow(img, cmap=plt.get_cmap('gray'))
img = np.asarray(img)
img = cv2.resize(img, (32, 32))
img = preprocessing(img)
plt.imshow(img, cmap = plt.get_cmap('gray'))
print(img.shape)
```




```
img = img.reshape(1, 32, 32, 1)
print("Predicted sign: " + str(np.argmax(model.predict(img), axis=-1)))
```

Predicted sign: [23]

```
r = requests.get(url14, stream=True)
img = Image.open(r.raw)
plt.imshow(img, cmap=plt.get_cmap('gray'))
img = np.asarray(img)
img = cv2.resize(img, (32, 32))
img = preprocessing(img)
plt.imshow(img, cmap = plt.get_cmap('gray'))
print(img.shape)
```


(32, 32)

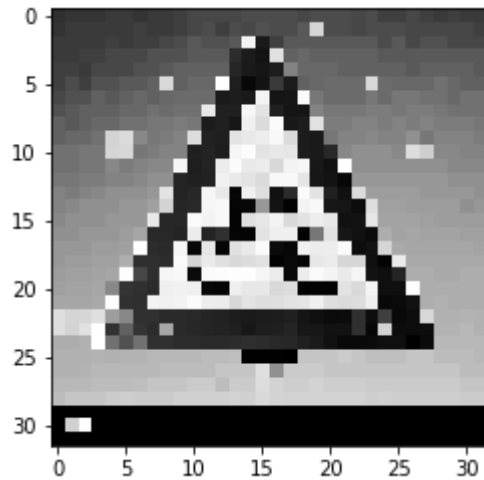


```
img = img.reshape(1, 32, 32, 1)
print("Predicted sign: "+ str(np.argmax(model.predict(img), axis=-1)))
```

Predicted sign: [13]

```
r = requests.get(url5, stream=True)
img = Image.open(r.raw)
plt.imshow(img, cmap=plt.get_cmap('gray'))
img = np.asarray(img)
img = cv2.resize(img, (32, 32))
img = preprocessing(img)
plt.imshow(img, cmap = plt.get_cmap('gray'))
print(img.shape)
```

(32, 32)



```
img = img.reshape(1, 32, 32, 1)
print("Predicted sign: "+ str(np.argmax(model.predict(img), axis=-1)))
```

```
Predicted sign: [29]
```

✓ 0s completed at 3:12 AM

